

INFS2200 Assignment, Semester 2 – 2017

Details

- **Student Number:** 43926871
- **Student Name:** Maxwell Bo
- **DOG_ID:** 988

Task 1: Constraints

a)

```
SQL> SELECT OWNER, CONSTRAINT_NAME, TABLE_NAME, SEARCH_CONDITION, INDEX_NAME FROM  
USER_CONSTRAINTS;
```

OWNER INDEX_NAME	CONSTRAINT_NAME	TABLE_NAME	SEARCH_CONDITION
-----	-----	-----	-----
C##S4392687 PK_STORES	PK_STORES	STORES	
C##S4392687	PK_DOG_BREEDS	DOG_BREEDS	

PK_DOG_BREEDS			
C##S4392687	PK_SERVICES	SERVICES	
PK_SERVICES			
C##S4392687	PK_DOGS	DOGS	
PK_DOGS			
C##S4392687	PK_SERVICE_HISTORY	SERVICE_HISTORY	
PK_SERVICE_HISTORY			
C##S4392687	PK_SHD	SERVICE_HISTORY_DETAIL	
PK_SHD			
C##S4392687	NN_PRICE	SERVICES	PRICE IS NOT NULL
C##S4392687	NN_DOG_NAME	DOGS	DOG_NAME IS NOT NULL
C##S4392687	FK_DOG_BREED	DOGS	
C##S4392687	FK_SH_STORE_ID	SERVICE_HISTORY	
C##S4392687	FK_SH_DOG_ID	SERVICE_HISTORY	
C##S4392687	FK_SHD_SERVICE_ID	SERVICE_HISTORY_DETAIL	
C##S4392687	FK_SERVICE_NAME	SERVICE_HISTORY_DETAIL	

13 rows selected.

b)

```
SQL> ALTER TABLE CUSTOMERS ADD CONSTRAINT "PK_CUSTOMERS" PRIMARY KEY (C_ID);
```

Table altered.

```
SQL> ALTER TABLE DOGS ADD CONSTRAINT "FK_C_ID" FOREIGN KEY (C_ID) REFERENCES  
CUSTOMERS(C_ID);
```

Table altered.

```
SQL> ALTER TABLE SERVICE_HISTORY_DETAIL ADD CONSTRAINT "FK_SHD_SERVICE_ID" FOREIGN  
KEY (SERVICE_ID) REFERENCES SERVICE_HISTORY(SERVICE_ID);
```

Table altered.

```
SQL> ALTER TABLE CUSTOMERS ADD CONSTRAINT "NN_DOB" CHECK (DOB IS NOT NULL);
```

Table altered.

```
SQL> ALTER TABLE SERVICE_HISTORY ADD CONSTRAINT "CK_FINISHED" CHECK (FINISHED IN ('T', 'F'));
```

Table altered.

```
SQL> ALTER TABLE CUSTOMERS ADD CONSTRAINT "CK_DOB" CHECK (DOB < DATE '1999-01-01');
```

Table altered.

```
SQL> ALTER TABLE SERVICE_HISTORY_DETAIL ADD CONSTRAINT "CK_START_TIME_END_TIME" CHECK (START_TIME < END_TIME);
```

Table altered.

```
SQL> ALTER TABLE SERVICE_HISTORY_DETAIL ADD CONSTRAINT "CK_SERVICE_DATE" CHECK (END_TIME < DATE '2018-01-01');
```

Table altered.

Task 2: Triggers

a)

```
SQL> CREATE OR REPLACE TRIGGER "TR_CUSTOMER_ID"  
2  BEFORE INSERT ON CUSTOMERS  
3  FOR EACH ROW  
4  BEGIN  
5  SELECT "SEQ_CUSTOMER".NEXTVAL INTO :NEW.C_ID FROM DUAL;  
6  END;  
7  /
```

Trigger created.

b)

```
SQL> CREATE OR REPLACE TRIGGER "TR_SERVICE_ID"  
  
2  BEFORE INSERT ON SERVICE_HISTORY  
  
3  FOR EACH ROW  
  
4  BEGIN  
  
5  SELECT SEQ_SERVICE_HISTORY.NEXTVAL INTO :NEW.SERVICE_ID FROM DUAL;  
  
6  END;  
  
7  /
```

Trigger created.

c)

```
SQL> CREATE OR REPLACE TRIGGER "TR_SERVICE_HISTORY_MESSAGE"

 2  BEFORE INSERT ON SERVICE_HISTORY

 3  FOR EACH ROW

 4  BEGIN

 5  IF :NEW.FINISHED = 'T' THEN

 6      SELECT 'Hi '

 7      || C.F_NAME || ' '

 8      || C.L_NAME || ', your dog '

 9      || D.DOG_NAME || ' of breed: '

10      || D.DOG_BREED || ' is ready for pick up at '

11      || S.STORE_AREA || '.'

12  INTO :NEW.MESSAGE

13  FROM CUSTOMERS C, DOGS D, STORES S
```

```
14 WHERE C.C_ID = D.C_ID

15 AND :NEW.STORE_ID = S.STORE_ID

16 AND :NEW.DOG_ID = D.DOG_ID;

17 ELSE

18 SELECT 'Hi '

19 || C.F_NAME || ' '

20 || C.L_NAME || ', your dog '

21 || D.DOG_NAME || ' of breed: '

22 || D.DOG_BREED|| ' is not ready to be picked up yet.'

23 INTO :NEW.MESSAGE

24 FROM CUSTOMERS C, DOGS D

25 WHERE C.C_ID = D.C_ID

26 AND :NEW.DOG_ID = D.DOG_ID;
```



```
27  END IF;
```

```
28  END;
```

```
29  /
```

Trigger created.

d)

```
SQL> INSERT INTO CUSTOMERS (F_NAME, L_NAME, DOB)
```

```
2  VALUES ('Luke', 'Cheung', '08-OCT-1996');
```

1 row created.

```
SQL> SELECT * FROM CUSTOMERS WHERE F_NAME='Luke' AND L_NAME='Cheung';
```

C_ID	F_NAME	L_NAME	DOB
10000	Luke	Cheung	08-OCT-96

```
SQL> INSERT INTO SERVICE_HISTORY (DOG_ID, STORE_ID, FINISHED)
```

```
2 VALUES (1234, 30, 'F');
```

```
1 row created.
```

```
SQL> SELECT * FROM SERVICE_HISTORY WHERE DOG_ID=1234 AND STORE_ID=30;
```

```
DOG_ID    STORE_ID SERVICE_ID F MESSAGE
```

```
-----  
1234      30      125000 F Hi Lady Finland, your dog Jack  
son of breed: English Foxhound  
is not ready to be picked up  
yet.
```

Task 3: Views

a)

```
SQL> CREATE VIEW "V_DOG_BREED_STATISTICS" AS  
2  SELECT D.DOG_BREED, SUM(S.PRICE) as TOTAL, AVG(S.PRICE) as MEAN,
```

```
STDDEV(S.PRICE) as STANDARD_DEVIATION
```

```
3 FROM DOGS D, SERVICE_HISTORY SH, SERVICE_HISTORY_DETAIL SHD, SERVICES S  
4 WHERE D.DOG_ID = SH.DOG_ID  
5 AND SH.SERVICE_ID = SHD.SERVICE_ID  
6 AND SHD.SERVICE_NAME = S.SERVICE_NAME  
7 GROUP BY D.DOG_BREED;
```

View created.

b)

```
SQL> CREATE MATERIALIZED VIEW "MV_DOG_BREED_STATISTICS"  
2 BUILD IMMEDIATE
```

```
3 AS

4 SELECT D.DOG_BREED, SUM(S.PRICE) as TOTAL, AVG(S.PRICE) as MEAN,
STDDEV(S.PRICE) as STANDARD_DEVIATION

5 FROM DOGS D, SERVICE_HISTORY SH, SERVICE_HISTORY_DETAIL SHD, SERVICES S

6 WHERE D.DOG_ID = SH.DOG_ID

7 AND SH.SERVICE_ID = SHD.SERVICE_ID

8 AND SHD.SERVICE_NAME = S.SERVICE_NAME

9 GROUP BY D.DOG_BREED;
```

Materialized view created.

c)

```
SQL> Select Command Prompt - sql
DOG_BREED
-----
STANDARD_DEVIATION
-----
Polish Lowland Sheepdog
      10.8171544
German Spaniel
      11.0714499
Saarlooswolfhond
      10.7575868

DOG_BREED
-----
STANDARD_DEVIATION
-----
Rat Terrier
      11.3098557
Drever
      10.9373436

515 rows selected.
Elapsed: 00:00:02.65
SQL>
```

```
SQL> Command Prompt - sqlplus
DOG_BREED
-----
STANDARD_DEVIATION
-----
Polish Lowland Sheepdog
      10.8171544
German Spaniel
      11.0714499
Saarlooswolfhond
      10.7575868

DOG_BREED
-----
STANDARD_DEVIATION
-----
Rat Terrier
      11.3098557
Drever
      10.9373436

515 rows selected.
Elapsed: 00:00:01.96
SQL>
```

A virtual view is simply another query masquerading as an in-memory table. When the virtual view was queried, the queries were composed together into a combined query, which was then executed.

When the materialized view was created, its query was immediately executed, and the result of the query was cached (forced by **BUILD IMMEDIATE**) . When the view was queried, the precomputed result was used instead, leading to a lower overall running time when compared to the virtual view.

Task 4: Function Based Indexes

a)

```
SQL> SELECT D.DOG_ID, D.DOG_NAME, DENTAL_CHECKUPS.DIFFERENCE, T.STORE_AREA
 2 FROM
 3 (SELECT SHD.SERVICE_ID, MAX(SHD.END_TIME - SHD.START_TIME) AS DIFFERENCE
 4 FROM SERVICE_HISTORY_DETAIL SHD, SERVICES S
 5 WHERE SHD.SERVICE_NAME = 'Dental Checkup'
 6 GROUP BY SHD.SERVICE_ID) DENTAL_CHECKUPS,
 7 SERVICE_HISTORY SH, STORES T, DOGS D
 8 WHERE DENTAL_CHECKUPS.SERVICE_ID = SH.SERVICE_ID
 9 AND SH.STORE_ID = T.STORE_ID
10 AND SH.DOG_ID = D.DOG_ID
11 ORDER BY DENTAL_CHECKUPS.DIFFERENCE DESC
12 FETCH FIRST 1 ROW ONLY;
```

DOG_ID	DOG_NAME	DIFFERENCE	STORE_AREA
5747	Layla	+000000000 01:18:00.000000	Bridgeman Downs

Elapsed: 00:00:01.27
SQL>

b)

```
SQL> CREATE INDEX "IDX_SERVICE_TIME" ON SERVICE_HISTORY_DETAIL(END_TIME -
START_TIME);
```

Index created.

c)

```
SP2-0042: unknown command "3ca" - rest of line ignored.
SQL> SELECT D.DOG_ID, D.DOG_NAME, DENTAL_CHECKUPS.DIFFERENCE, T.STORE_AREA
 2 FROM
 3 (SELECT SHD.SERVICE_ID, MAX(SHD.END_TIME - SHD.START_TIME) AS DIFFERENCE
 4 FROM SERVICE_HISTORY_DETAIL SHD, SERVICES S
 5 WHERE SHD.SERVICE_NAME = 'Dental Checkup'
 6 GROUP BY SHD.SERVICE_ID) DENTAL_CHECKUPS,
 7 SERVICE_HISTORY SH, STORES T, DOGS D
 8 WHERE DENTAL_CHECKUPS.SERVICE_ID = SH.SERVICE_ID
 9 AND SH.STORE_ID = T.STORE_ID
10 AND SH.DOG_ID = D.DOG_ID
11 ORDER BY DENTAL_CHECKUPS.DIFFERENCE DESC
12 FETCH FIRST 1 ROW ONLY;

   DOG_ID DOG_NAME
-----
DIFFERENCE
-----
STORE_AREA
-----
      5747 Layla
+0000000000 01:18:00.000000
Bridgeman Downs

Elapsed: 00:00:01.13
SQL>
```

The index has reduced the running time. A function based index essentially precomputes the given expression for each table row in the specified table. Queries that use this expression can refer to this index. This index must be maintained, similar to a materialized view.

Task 5: Bitmap Indexing

a)

```
SQL> SELECT SERVICE_NAME, COUNT(1)
  2  FROM SERVICE_HISTORY_DETAIL
  3  GROUP BY SERVICE_NAME;
```

SERVICE_NAME	COUNT(1)
Flea Prevention	46594
Fur Trim	46652
Breath Treatment	46742
Nail Trim	46551
Tapeworm Prevention	46807
Wash	46781
Paw Pad Treatment	46738
Dental Checkup	46421

8 rows selected.

Elapsed: 00:00:00.05

b)

```
SQL> CREATE BITMAP INDEX "BIDX_SERVICE" ON SERVICE_HISTORY_DETAIL(SERVICE_NAME);
```

Index created.

c)

```
SQL> SELECT SERVICE_NAME, COUNT(1)
  2  FROM SERVICE_HISTORY_DETAIL
  3  GROUP BY SERVICE_NAME;
```

SERVICE_NAME	COUNT(1)
Breath Treatment	46742
Dental Checkup	46421
Flea Prevention	46594
Fur Trim	46652
Nail Trim	46551
Paw Pad Treatment	46738
Tapeworm Prevention	46807
Wash	46781

8 rows selected.

Elapsed: 00:00:00.02

The index has reduced the running time. Before adding the bitmap index, the query needed to go through each row, and check the **SERVICE_NAME** with either a hash function or a string comparison in order to appropriately tally each unique **SERVICE_NAME**. After adding the bitmap index, the tally can be quickly generated by using a aggregation function that sweeps the bitmap. The bitmap has a column for each unique

SERVICE_NAME, and indicates the row's **SERVICE_NAME** by setting a bit in the appropriate column.

d)

The advantages of constructing this bitmap index include dramatically increased performance with queries that perform aggregation.

The disadvantages of constructing this bitmap index include having to reconfigure the bitmap every time a new unique value is introduced to a column, and maintaining the map when altering pre-existing values.

TASK 6: Execution Plan & Analysis

a)

```
SQL> SELECT
  2  GET_UNIQUE_SNUMBER(43926871) FROM
  3  DUAL;
```

```
GET_UNIQUE_SNUMBER(43926871)
-----
                              988
```

```
Elapsed: 00:00:00.01
```

b)

```
SQL> SELECT INDEX_NAME, COLUMN_NAME, TABLE_NAME
2  FROM USER_IND_COLUMNS
3  WHERE TABLE_NAME='STORES' OR TABLE_NAME='SERVICE_HISTORY' OR
TABLE_NAME='SERVICE_HISTORY_DETAIL';
```

INDEX_NAME	COLUMN_NAME	TABLE_NAME
-----	-----	-----
PK_SERVICE_HISTORY	SERVICE_ID	SERVICE_HISTORY
PK_SHD	SERVICE_ID	SERVICE_HISTORY_DETAIL
PK_SHD	SERVICE_NAME	SERVICE_HISTORY_DETAIL
BIDX_SERVICE	SERVICE_NAME	SERVICE_HISTORY_DETAIL
IDX_SERVICE_TIME	SYS_NC00005\$	SERVICE_HISTORY_DETAIL
PK_STORES	STORE_ID	STORES

6 rows selected.

c)

```
SQL> SELECT COUNT(*)
  2  FROM SERVICE_HISTORY SH, SERVICE_HISTORY_DETAIL SHD
  3  WHERE SH.DOG_ID = 988
  4  AND SH.SERVICE_ID = SHD.SERVICE_ID;

COUNT(*)
-----
         45

Elapsed: 00:00:00.14
```

```
SQL> EXPLAIN PLAN FOR (SELECT COUNT(*)
  2  FROM SERVICE_HISTORY SH, SERVICE_HISTORY_DETAIL SHD
  3  WHERE SH.DOG_ID = 988
  4  AND SH.SERVICE_ID = SHD.SERVICE_ID);
```

Explained.

```
SQL> SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

```
PLAN_TABLE_OUTPUT
-----
```

```
Plan hash value: 2851704051
```


Id	Operation	Name	Rows	Bytes	Cost
(%CPU)	Time				

0	SELECT STATEMENT		1	39	556
(2)	00:00:01				
1	SORT AGGREGATE		1	39	
2	NESTED LOOPS		2450	95550	556
(2)	00:00:01				
3	NESTED LOOPS		327K	95550	556
(2)	00:00:01				
4	INDEX FAST FULL SCAN	PK_SHD	327K	4159K	545
(0)	00:00:01				
* 5	INDEX UNIQUE SCAN	PK_SERVICE_HISTORY	1		0
(0)	00:00:01				
* 6	TABLE ACCESS BY INDEX ROWID	SERVICE_HISTORY	1	26	0
(0)	00:00:01				

Predicate Information (identified by operation id):

```
5 - access("SH"."SERVICE_ID"="SHD"."SERVICE_ID")
6 - filter("SH"."DOG_ID"=988)
```

Note

- dynamic statistics used: dynamic sampling (level=2)

23 rows selected.

The plan is executed in the following order:

- **INDEX UNIQUE SCAN.** Index **PK_SERVICE_HISTORY** is used in a unique scan operation to evaluate the **WHERE** clause criteria. It returns exactly one rowid from the index, and is a traversal through a B+ tree.
- **TABLE ACCESS BY INDEX ROWID.** Rows are located using index(es), iff the table is nonpartitioned.
- **INDEX FAST FULL SCAN.** Retrieval of all rowids from an index when there is no start or stop key. Indexed values are scanned in ascending order.
- **SORT AGGREGATE.** Retrieval of a single row that is the result of applying a group function to a group of selected rows. In this instance, the group function is **COUNT**.
- **SELECT STATEMENT.** Returns rows satisfying the **WHERE** clause conditions.

d)

```
SQL> ALTER TABLE SERVICE_HISTORY_DETAIL DROP CONSTRAINT PK_SHD;
```

Table altered.

```
SQL> ALTER TABLE SERVICE_HISTORY_DETAIL DROP CONSTRAINT PK_SHD;
```

Table altered.

```
SQL> ALTER TABLE SERVICE_HISTORY DROP CONSTRAINT PK_SERVICE_HISTORY;
```

Table altered.


```
SQL> SELECT COUNT(*)
  2  FROM SERVICE_HISTORY SH, SERVICE_HISTORY_DETAIL SHD
  3  WHERE SH.DOG_ID = 988
  4  AND SH.SERVICE_ID = SHD.SERVICE_ID;

COUNT(*)
-----
         45

Elapsed: 00:00:00.23
```

```
SQL> EXPLAIN PLAN FOR (SELECT COUNT(*)
  2  FROM SERVICE_HISTORY SH, SERVICE_HISTORY_DETAIL SHD
  3  WHERE SH.DOG_ID = 988
  4  AND SH.SERVICE_ID = SHD.SERVICE_ID);
```

Explained.

```
SQL> SELECT PLAN_TABLE_OUTPUT FROM TABLE (DBMS_XPLAN.DISPLAY);
```

```
PLAN_TABLE_OUTPUT
-----
```

```
Plan hash value: 3313643833
```

```

-----
| Id | Operation | Name | Rows | Bytes | Cost
(%CPU)| Time |
-----
| 0 | SELECT STATEMENT | | 1 | 39 | 1198
(1)| 00:00:01 |
| 1 | SORT AGGREGATE | | 1 | 39 |
|
|* 2 | HASH JOIN | | 2450 | 95550 | 1198
(1)| 00:00:01 |
|* 3 | TABLE ACCESS FULL | SERVICE_HISTORY | 25 | 650 | 615
(1)| 00:00:01 |
| 4 | TABLE ACCESS FULL | SERVICE_HISTORY_DETAIL | 327K | 4159K | 582
(1)| 00:00:01 |
-----

```

Predicate Information (identified by operation id):

```

2 - access("SH"."SERVICE_ID"="SHD"."SERVICE_ID")
3 - filter("SH"."DOG_ID"=988)

```

Note

- dynamic statistics used: dynamic sampling (level=2)

21 rows selected.

The plan is executed in the following order.

- **TABLE ACCESS FULL**. Retrieval of all rows from the **SERVICE_HISTORY_DETAIL** table.
- **TABLE ACCESS FULL**. Retrieval of all rows from the **SERVICE_HISTORY** table.
- **HASH JOIN**. Operation joining two sets of rows and returning the result. The join condition (**SERVICE_ID**) is an efficient way of accessing the second table.
- **SORT AGGREGATE**. Retrieval of a single row that is the result of applying a group function to a group of selected rows. In this instance, the group function is **COUNT**.
- **SELECT STATEMENT**. Returns rows satisfying the **WHERE** clause conditions.

e)

```
SQL> SELECT HEIGHT, LF_BLKs, BLKS_GETS_PER_ACCESS FROM INDEX_STATS;
```

HEIGHT	LF_BLKs	BLKS_GETS_PER_ACCESS
--------	---------	----------------------

1

1

2

The height of the tree is 1. There is 1 leaf block. Only 2 block accesses are needed for a full direct scan of the **STORES** table.